

22nd April Assignment

April 29, 2023

1 Assignment 75

Q1. Write a Python code to implement the KNN classifier algorithm on load_iris dataset in sklearn.datasets.

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.datasets import load_iris
```

```
[2]: dataset = load_iris()
```

```
[3]: df = pd.DataFrame(data=dataset.data, columns=dataset.feature_names)
df['Target'] = dataset.target
df.head()
```

```
[3]:      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0                5.1             3.5             1.4             0.2
1                4.9             3.0             1.4             0.2
2                4.7             3.2             1.3             0.2
3                4.6             3.1             1.5             0.2
4                5.0             3.6             1.4             0.2
```

```
      Target
0         0
1         0
2         0
3         0
4         0
```

```
[4]: df.isnull().sum()
```

```
[4]: sepal length (cm)    0
sepal width (cm)       0
petal length (cm)      0
petal width (cm)       0
```

```
Target          0
dtype: int64
```

```
[5]: # segregate the data into independent and dependent features
X = df.iloc[:, :-1]
y = df['Target']
```

```
[6]: # train test and split
from sklearn.model_selection import train_test_split
```

```
[8]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↳20,random_state=42)
```

```
[9]: # model training
from sklearn.neighbors import KNeighborsClassifier
```

```
[10]: classifier = KNeighborsClassifier(n_neighbors=5,algorithm='auto')
```

```
[11]: classifier.fit(X_train,y_train)
```

```
[11]: KNeighborsClassifier()
```

```
[12]: y_pred = classifier.predict(X_test)
```

```
[13]: # Evaluation of performance
from sklearn.metrics import confusion_matrix, accuracy_score,
↳classification_report
```

```
[14]: print(confusion_matrix(y_pred,y_test))
print(accuracy_score(y_pred,y_test))
print(classification_report(y_pred,y_test))
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

```
1.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Q2. Write a Python code to implement the KNN regressor algorithm on load_boston dataset in sklearn.datasets.

```
[15]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import urllib.request
```

```
[16]: # Download the dataset from its URL
url = 'https://raw.githubusercontent.com/selva86/datasets/master/BostonHousing.
      ↪csv'
urllib.request.urlretrieve(url, 'BostonHousing.csv')

# Load the dataset into a Pandas DataFrame
df = pd.read_csv('BostonHousing.csv')
```

```
[17]: df.head()
```

```
[17]:      crim    zn  indus  chas   nox    rm   age   dis  rad  tax  ptratio  \
0  0.00632  18.0   2.31    0  0.538  6.575  65.2  4.0900   1  296    15.3
1  0.02731   0.0   7.07    0  0.469  6.421  78.9  4.9671   2  242    17.8
2  0.02729   0.0   7.07    0  0.469  7.185  61.1  4.9671   2  242    17.8
3  0.03237   0.0   2.18    0  0.458  6.998  45.8  6.0622   3  222    18.7
4  0.06905   0.0   2.18    0  0.458  7.147  54.2  6.0622   3  222    18.7

      b  lstat  medv
0  396.90   4.98  24.0
1  396.90   9.14  21.6
2  392.83   4.03  34.7
3  394.63   2.94  33.4
4  396.90   5.33  36.2
```

```
[18]: df.isnull().sum()
```

```
[18]: crim      0
      zn        0
      indus     0
      chas      0
      nox       0
      rm        0
      age       0
      dis       0
      rad       0
      tax       0
      ptratio   0
      b         0
```

```
lstat      0
medv       0
dtype: int64
```

```
[19]: # segregate the feature into independent and dependent feature
X = df.iloc[:, :-1]
y = df['medv']
```

```
[20]: # train test and split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↳20, random_state=42)
```

```
[21]: # model training
from sklearn.neighbors import KNeighborsRegressor
```

```
[22]: reg = KNeighborsRegressor()
```

```
[23]: reg.fit(X_train, y_train)
```

```
[23]: KNeighborsRegressor()
```

```
[24]: y_pred = reg.predict(X_test)
```

```
[25]: # Evaluation of performance
from sklearn.metrics import r2_score
```

```
[26]: print(r2_score(y_test, y_pred))
```

```
0.6473640882039258
```

Q3. Write a Python code snippet to find the optimal value of K for the KNN classifier algorithm using cross-validation on load_iris dataset in sklearn.datasets.

```
[27]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.datasets import load_iris
```

```
[28]: dataset = load_iris()
```

```
[29]: df = pd.DataFrame(data=dataset.data, columns=dataset.feature_names)
df['Target'] = dataset.target
df.head()
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	Target
0	0
1	0
2	0
3	0
4	0

```
[30]: # segregate the data into independent and dependent features
X = df.iloc[:, :-1]
y = df['Target']
```

```
[31]: # train test and split
from sklearn.model_selection import train_test_split
```

```
[32]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.
↳20,random_state=42)
```

```
[33]: # model training
from sklearn.neighbors import KNeighborsClassifier
```

```
[34]: classifier = KNeighborsClassifier()
```

```
[35]: # Task
from sklearn.model_selection import GridSearchCV
K_val = range(1,11)
parameters = {
    'n_neighbors':K_val
}
```

```
[36]: # Hypertunning
from sklearn.model_selection import GridSearchCV
```

```
[37]: grid = GridSearchCV(classifier,param_grid=parameters,scoring='accuracy')
```

```
[38]: grid.fit(X_train,y_train)
```

```
[38]: GridSearchCV(estimator=KNeighborsClassifier(),
    param_grid={'n_neighbors': range(1, 11)}, scoring='accuracy')
```

```
[40]: grid.best_params_
```

```
[40]: {'n_neighbors': 3}
```

```
[41]: classifier= KNeighborsClassifier(**grid.best_params_)
```

```
[42]: classifier.fit(X_train,y_train)
```

```
[42]: KNeighborsClassifier(n_neighbors=3)
```

```
[43]: y_pred = classifier.predict(X_test)
```

```
[44]: # Evaluation of performance
from sklearn.metrics import confusion_matrix, accuracy_score, \
    classification_report
print(confusion_matrix(y_pred,y_test))
print(accuracy_score(y_pred,y_test))
print(classification_report(y_pred,y_test))
```

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
1.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Q4. Implement the KNN regressor algorithm with feature scaling on load_boston dataset in sklearn.datasets.

```
[45]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import urllib.request
```

```
[46]: # Download the dataset from its URL
url = 'https://raw.githubusercontent.com/selva86/datasets/master/BostonHousing.
    ↪CSV'
urllib.request.urlretrieve(url, 'BostonHousing.csv')

# Load the dataset into a Pandas DataFrame
```

```
df = pd.read_csv('BostonHousing.csv')
```

```
[47]: df.head()
```

```
[47]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

```
[48]: # segregate the feature into independent and dependent feature
X = df.iloc[:, :-1]
y = df['medv']
```

```
[49]: # train test and split
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
↳ 20, random_state=42)
```

```
[50]: # feature scaling
from sklearn.preprocessing import StandardScaler
```

```
[51]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
[52]: # model training
from sklearn.neighbors import KNeighborsRegressor
```

```
[53]: # Train the KNN regressor
knn_regressor = KNeighborsRegressor(n_neighbors=5)
knn_regressor.fit(X_train_scaled, y_train)
```

```
[53]: KNeighborsRegressor()
```

```
[54]: # Make predictions on the test set
y_pred = knn_regressor.predict(X_test_scaled)
```

```
[55]: # Evaluation of performance
from sklearn.metrics import mean_squared_error, r2_score
print(mean_squared_error(y_test,y_pred))
print(r2_score(y_test,y_pred))
```

20.60552941176471
0.7190172315709293

Q5. Write a Python code snippet to implement the KNN classifier algorithm with weighted voting on load_iris dataset in sklearn.datasets.

```
[56]: from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

# Load the Iris dataset
X, y = load_iris(return_X_y=True)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    random_state=42)

# Train the KNN classifier with weighted voting
knn_classifier = KNeighborsClassifier(n_neighbors=5, weights='distance')
knn_classifier.fit(X_train, y_train)

# Make predictions on the test set
y_pred = knn_classifier.predict(X_test)

# Evaluate the model using accuracy score
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```

Accuracy: 1.0

Q6. Implement a function to standardise the features before applying KNN classifier.

```
[57]: from sklearn.preprocessing import StandardScaler

def knn_classifier(X_train, X_test, y_train, y_test, n_neighbors):
    # Standardize the features
    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    # Train the KNN classifier
    knn = KNeighborsClassifier(n_neighbors=n_neighbors)
```



```

knn.fit(X_train_scaled, y_train)

# Make predictions on the test set
y_pred = knn.predict(X_test_scaled)

# Return the predictions and the trained model
return y_pred, knn

```

Q7. Write a Python function to calculate the euclidean distance between two points.

```

[58]: import numpy as np
def euclidean_distance(point1, point2):
    point1 = np.asarray(point1)
    point2 = np.asarray(point2)
    distance = np.sqrt(np.sum((point1 - point2) ** 2))
    return distance

```

```

[59]: point1 = [1, 2, 3]
point2 = [4, 5, 6]
distance = euclidean_distance(point1, point2)
print(distance)

```

5.196152422706632

Q8. Write a Python function to calculate the manhattan distance between two points.

```

[60]: import numpy as np

def manhattan_distance(point1, point2):
    point1 = np.asarray(point1)
    point2 = np.asarray(point2)
    distance = np.sum(np.abs(point1 - point2))
    return distance

```

```

[61]: point1 = [1, 2, 3]
point2 = [4, 5, 6]
distance = manhattan_distance(point1, point2)
print(distance)

```